Taxonomic classification of ants (Formicidae) from images using deep learning

Marijn J. A. Boer [1,2,*]

Abstract

The diverse and species rich group of the ants (Formicidae) consists of important biodindicators for ecosystem health and could be surrogates for overall species richness, but identifying ant species is hard and requires specialists. In the field of computer vision, machine learning has shown promising results in species identification and could greatly speed up and improve taxonomic accuracy. While previous work has been conventional in that features needed to be extracted (i.e. supervised learning), here, deep learning (in the form of a convolutional neural network) is proposed to learn features automatically in order to classify ant species. Such a system could answer whether ant species could be discriminated using deep learning, insight that will contribute to Naturalis's broader goals for computer vision. The use of ant images will also raise questions on the use of different castes and different image shot types, which will be combined in a multi-view approach. For this prototype, a dataset with the 97 most imaged species (10,204 images) from AntWeb was used. The dataset shows little difference in shot type, but there is a large distribution in the width of the image dimension. As part of the prototype, two options for training and validating the classifier are explored: (1) use of an existing, already promising CNN and (2) designing a CNN by start from a simple template that is expanded after each evaluation. The trained networks are evaluated using a top-1 and top-5 accuracy test with the aim to classify morphospecies to valid species.

Keywords

Convolutional Neural Network — Image Classification — AntWeb — Ants — Formicidae

1 *Rutger Vos Lab, Endless Forms, Naturalis Biodiversity Center, Leiden, Pesthuislaan 7, 2333 BA, Leiden, Netherlands*
2 *Masterstudent Biology; Evolution, Conservation and Biodiversity, Institute of Biology Leiden, Sylviusweg 72, 2333 BE, Leiden, Netherlands*
*** Corresponding author**: marijn.boer@naturalis.nl

1. Introduction

1.1 Taxonomy and the ants

The family of the ants (Formicidae) is a large and diverse group within the insect order, sometimes exceeding other groups of insects in local diversity by far. They are globally found (with the exception of Antarctica) and play an important role in a lot of ecosystems [1]. To get a good understanding of their role in ecosystems, it is important to have information on their species richness and other indicators of diversity. With this information, ants can also act as a biopredictor for the health of ecosystems, which is in turn important for species conservation [2]. Furthermore, arthropods are found to be good surrogates for predicting species richness patterns in other groups of animals [3]. This may even be possible using morphospecies as taxa [4], but for that one still needs a species concept.

Understanding which species there are, is complex and specialists are generally required for correct identifications. Ant taxonomists use distinct characters, sometimes varying for genera and species, to perform identifications. Examples of these characters are antennae, hairs, carinae, thorax shape, shininess of the body. And still, it is difficult to identify species correctly, because, among other reasons, some ant species are very cryptic or sibling species, species have different castes and intercastes and male individuals are underrepresented in descriptions and collecting. Here, we explore an alternative approach to taxonomic identification based on computer vision and deep learning.

1.2 Computer Vision

Research in the field of computer vision shows prior art on insect identification from images. Wang et al. [5] designed an identification system to identify insects at the order level, including ants. They used 7 geometrical features in their system and accuracy reached 97%. In other research, the wing outline is used for insect identification [6]. Automating identification could greatly speed up taxonomic work and improve accuracy [7].

Automatic identification on species level has been subject of research for a while now. Machine

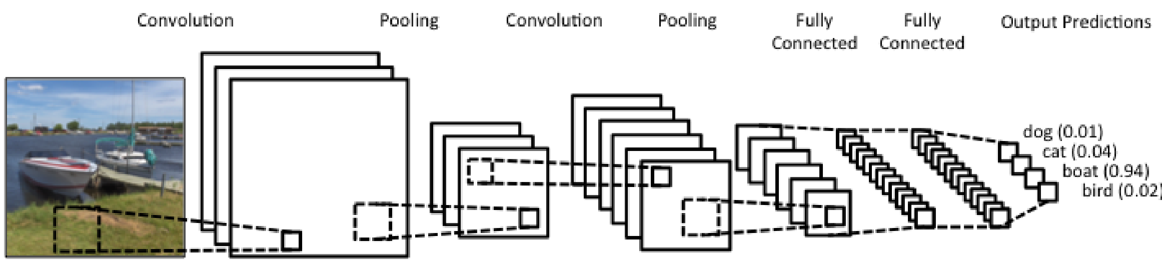

Figure 1. A simple representation of a Convolutional Neural Network. Source: clarifai.com/static/images/cnn.png

learning (ML) is applied on morphometrics of butterflies [8] with success. This was also done, using wing morphometrics of honey bees [9, 10]. And it is not limited to insects or even animals, as researchers designed a system to automatically identify 184 American trees using leaf morphometrics and color distributions with great success [11]. Others found that it was possible to use ML for classifying orchid species [12, 13]. Almost all of previous research exploit metrical features and morphometrics, which therefore still needs human supervision.

1.3 Deep learning

Deep learning (DL) shows great promise for taxonomic identification without supervision. In the past 5 to 7 years, DL attracted lots of attention in research and methods and algorithms have greatly improved [14]. DL allows a machine to learn features by itself, instead of conventional ways where features need to be introduced to the machine [14, 15]. DL has shown to perform much better when looking at multiple features, instead of a single one, as in conventional ways [16]. A successful DL algorithm is the convolutional neural network (CNN), mostly used for image classification and preferably trained using GPUs. These computationally intensive networks are designed to process (convolve) 2D data (mainly images), using typical layers as convolutional and pooling layers [14]. An example representation of a CNN can be seen in Figure 1 on page 2. In recent years, CNNs have advanced a lot [17] and different studies have shown promising results. CNNs have been used in plant identification identification [18, 19], plant disease detection [20] and even identification of underwater fish images [21], all with promising results. A more practical example is the classification of different qualities of wood for industrial purposes [22]. However, as promising as it sounds, DL still has some (mostly non-programming) obstacles to overcome [23].

1.4 Paving the way

Here, we explore the possibilities for automatic insect identification using DL. Because of rapid improvements in deep learning and computer power, a convolutional neural network is proposed as a feasible approach for automatic species identification. Our research uses a difficult dataset, ants (in contrast e.g. to colorful and pattern-rich butterflies), for which promising results should pave the way towards more general application in insect classification. Furthermore, the research should benefit Naturalis's goal of initializing an image classification system for own use. To explore CNNs, a trial is performed to incorporate a "three type view per specimen" approach. Almost all imaged specimens on AntWeb follow the same procedure of head, dorsal and profile shot type images. The combination of these three images per specimen should produce higher accuracy than when using only one image, as there is more data to train on. Because this is a fairly new and undiscovered field, with relatively little research published so far [24, 25].

Our main research question is quite simply whether a morphologically complex group such as ants can be discriminated from images using deep learning?. To answer this question, subquestions need to be answered while exploring subsets of the data. Three important subquestion will be (1) whether three shot types classify species more accurately than a single shot view?, (2) which shot type classifies species more accurately? and (3) which image quality has the best accuracy vs. training speed?. In relation to the image quality there is also a image dimension problem, as not all images have the same width.

In addition, there is the problem of different castes: what is the effect of using males, queens or workers; alate (winged) or dealate (non-winged) queens; major and minor workers; ergatoid and intercastes specimens? Ergatoids (productive) and inter-castes (non-productive) denotes the wingless individuals that look like the intermediate of a queen or male and a worker [26]. Also, what happens when morphospecies are incorporated, as they could be viewed as real species. Once the network is running with fair accuracy, how will it handle *indet* species? Most of these problems can be overcome, but others may not be. Lastly, AntWeb has amber resin and fossil specimens, specimens with broken or lost body parts (e.g. antennae), and even some incorrect identifications. For example, there is debate on North American and European *Lasius spp.* [27].

Programming code and documentation is available for open access (MIT licensed) and published on URL: `github.com/naturalis/FormicID`.

2. Proposed Approach & Methods

Our trials are on a relatively small dataset with a simple neural network to make the project more scalable and research more comprehensible. After evaluation of these results, new elements can be added to increase complexity and answer more complex questions. Each time after training where success is reached, extra dimensions of training data may be added. The following points line up examples of additions for increasing the complexity of this research:

- Use RGB images instead of grayscale
- Increase in image quality from low quality to medium quality
- Using all the valid species on AntWeb
- Including morphospecies
- Including alate queens and/or dealate queens and/or males and/or ergatoid specimens
- combining head, profile and dorsal shot views as a multi-view experiment

2.1 AntWeb

The data for this research is administered on the web page of AntWeb (URL: `antweb.org`), a website that provides a considerable amount of information on the diverse group of ants (Formicidae) [28]. The website has an extensive and well documented database of ant records. Specimens are, among other things, taxonomically classified, ecologically documented and imaged, mostly with sufficient accuracy. Data from AntWeb can be accessed using the website's API (version 2 [29] or version 3 beta [30]). Using programming scripts, the catalog number, scientific name, shot type and the images for a specimen are harvested. AntWeb's API version 2 does not support filtering for different castes when harvesting information, but with the updated version (version: 3 beta) it should be possible. At the moment version 2 is used, so all castes are represented when scraping images, which could give rise to problems. Furthermore, images on AntWeb are, for a significant part, all divided in 3 shot types per specimen; head, dorsal and profile views. Because of server restrictions images have to be downloaded in batches. AntWeb provides 4 different qualities of images, from lowest to highest; low, medium, thumbview, high. A representation of a species can be seen in Figure 2 on page 3.

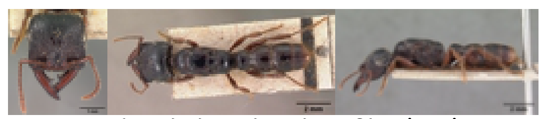

Figure 2. A head, dorsal and profile view image of *Amblyopone australis*. Images are in color and downloaded at the lowest quality.

2.2 Data preparation and processing

The first dataset for this research needed to be comprehensible, so a start can be made on the architecture of the neural network and fine-tuning all the algorithms. Images need to be in a grayscale format, evenly sized, correctly named and organized. Below a description is given on the first created dataset for this research, in order to compute a neural network and get preliminary results, before adding extra information to increase complexity.

2.2.1 *FormicID-97*

This first dataset, called *FormicID-97*, was made to have about 100 species (or classifiers). After ranking the most imaged species (combined for head, profile and dorsal views) a cut was made around a 100 species for a number of images. At this cut, the lowest ranking species had 68 images and the cut needed to be made at 101 species. This accounted for 10,225 AntWeb images. Out of these 101 species, 4 species were actually not representing real species, but the aggregated records of specimens that could not be identified to species in that genus. These groups were *Camponotus indet, Cremagotaster indet, Pheidole indet* and *Polyrhachis* *indet* and they were removed from the dataset, because it is unknown to see which species these specimens represent and even if they are the same species. Furthermore, a total of 21 specimens did not have valid URL's to images therefore could not be downloaded automatically. The dataset ended up with 97 species and a total of 10,204 images, which were downloaded to a directory and correctly named. All images were downloaded in color, at the lowest quality available. Images will be converted to grayscale speed up training in the first few trials.

The distribution of these species is shown in Figure 3 on page 5. This dataset is made up of head, profile and dorsal shot images of these species, but this is not evenly distributed as can be seen in Figure 4 on page 6. The dataset still needs to be filtered on castes, to get a worker only dataset.

Image height is at a constant 122 pixels for all images, but image width varies a lot as can be seen in Figure 5 on page 6. This problem could be potentially be solved using 1 or more of three methods. The first option is to squash the image to a square and thereby have all the images in the same dimensions. This would however change the shape of species, which could be an important classifying factor. The next option is to keep the aspect ratio, but to crop the image to the center using a set height and width. But this could erase parts of the body. The last option is to add black borders (i.e. zero's) on the image and then crop to get the desired height and width. The problem here is that this would add noise to the dataset, which the network would need to learn to ignore. The choice for one of these methods has not been made yet.

The file sizes are distributed rather normally around are around 20 kilobytes (Figure 6 on page 6). At the moment of writing all *FormicID-97* images are downloaded. The dataset will to be cleaned of non-workers castes and of problematic images. Then all images will be converted to grayscale images and resized to all have the exact same size. In order to increase the dataset, image augmentation can be applied, meaning that copies of images will be shifted, rotated and flipped vertically and added to the dataset.

In the first experiment *FormicID-97* will be split in three groups: head, profile and dorsal shots. This will ensure that the network is only training on a specific shot type and can not get confused by multiple views of a species. Results will show which shot type will better classify an image at species level. At this moment the dataset contains all types of castes; workers, alate and dealate queens, males and possible ergatoid or intercaste specimens. The combination of workers, queens and males will probably already make classification more complicated, but adding these intermediate forms will increase this even more.

For training purposes the 3 shot type-datasets will be split in training, validation and test sets, respectively in an 8:1:1 ratio, while keeping the species ratios intact.

2.3 Network architecture

All programming will be written in Python 3.6. For deep learning Keras will be used[31], with Tensorflow as backend [32]. There are two options in using a convolutional neural network (CNN).

1. The first option is to use an existing and promising CNN. There are multiple candidates for this and they all try to better the other. Below is a list of candidate CNNs:
    - *Resnet* [33]
    - *Inception-v4* / *Inception-ResNet* [34],
    - *DenseNet* [35]
    - *Xception* [36]
    - *GoogleNet* [37]
    - *VGGNet* [38]
2. They other option is to build a new neural network. This networks architecture will start as a relatively simple 3 layer network.

Furthermore, when designing a network there are a lot of things one need to think about. Below is a list of design choices that are made for this research. For now these are with reservation.

- Programming the network in Keras [31] using Python 3.6
- Use dropout [39]
- Choose NADAM as optimizers[40] (other possible optimizers are ADAM [41] or SGD [42, 43]
- Use Batch normalization [44]
- Perofrm parameter optimization [45]
- Visualizing the CNN architecture [46, 47, 48]
- Validation of network using different methdos (e.g. KNN or k-Fold cross-validation)

2.4 Testing network

The network will be tested for best top-5 accuracy and top-1 accuracy using the trained network. If accuracy is high *indet* species could be tested to see which species they could represent.

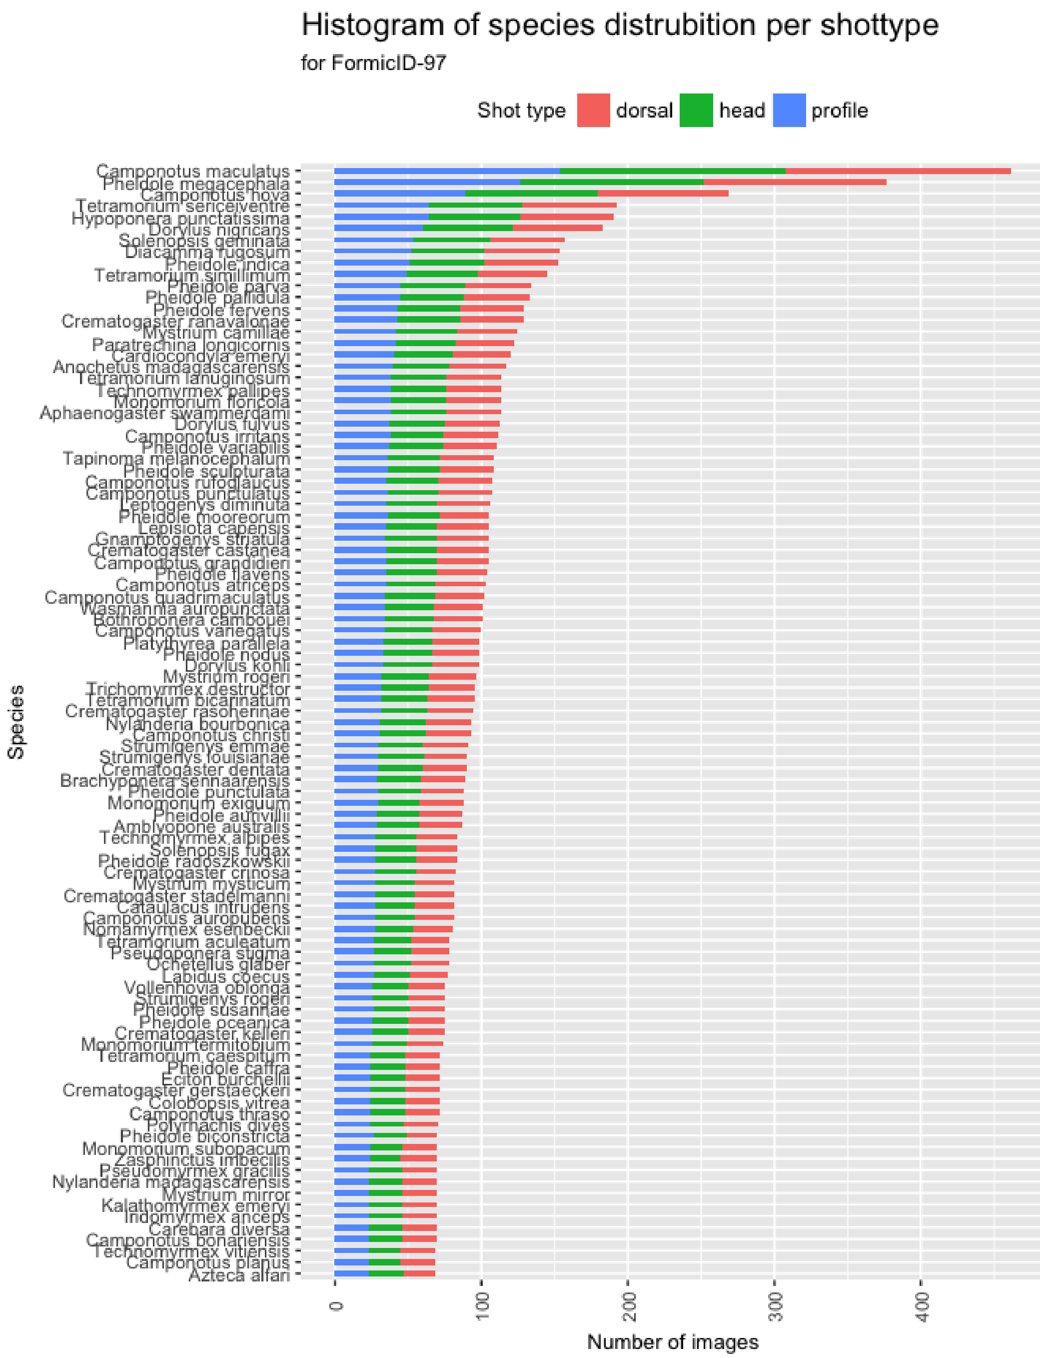

Figure 3. Histogram showing the distribution for the 97 most imaged species. Images are divided in dorsal (red), head (green) and profile (blue) views. Distribution ranges from the most imaged *Camponotus maculatus* with 462 images to *Azteca alfari, Camponotus planus* and *Technomyrmex vitiensis* with 68 images.

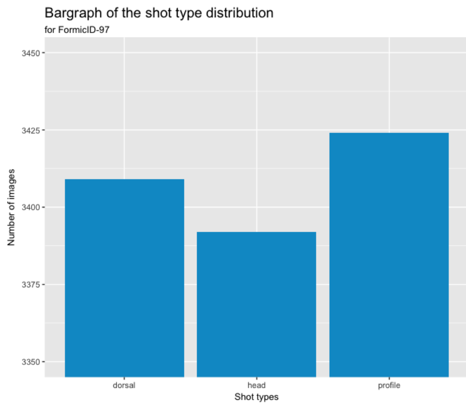

Figure 4. Bargraph showing the distribution for different shot types for the *FormicID-97* dataset. The y-axis is zoomed in from 3350 to 3450 in order to see a difference, but in fact the difference is very small.

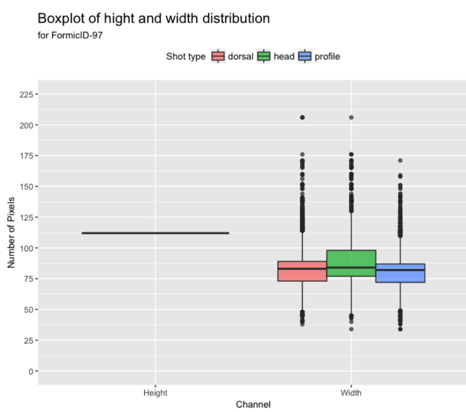

Figure 5. Boxplot showing the distribution for the number of height and width pixels for the *FormicID-97* dataset. Height is the same for all image, at exact 122 pixels.

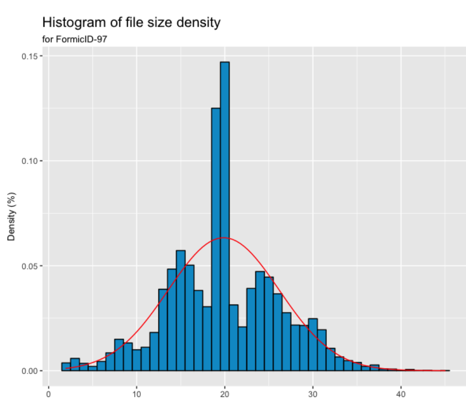

Figure 6. Distribution of the file size of the file sizes of all the RGB images from *FormicID-97*. The mean is 19.88 Kb.

2.5 Timeline

Figure 7 on page 7 shows the time schedule for this internship. The project start was on September 25th, 2017 and the deadline is set on July 31th, 2018. The proposal deadline is on is February 6th, 2018, accompanied by the proposal presentation. The end presentation will be planned later on. In December, 2017 the project was on a pause for 6 weeks for the purpose of education (course: Methods in Biodiversity Analysis - Leiden University) and a 2 week vacation. The timeline shows different parts of this project per week and in some weeks, subjects overlap, meaning that work will be combined or split in that week.

Acknowledgments

Lars Hopman and Joost Martens are thanked for helping the author as a beginning Python programmer and directions on starting with machine learning.

References

[1] B. Hölldobler and E.O. Wilson. *The Ants*. English. Cambridge, Mass.: Belknap Press of Harvard University Press, 1990. ISBN: 0-674-04075-9.

[2] Alan N. Andersen et al. "Using ants as bioindicators in land management: simplifying assessment of ant community responses". In: *Journal of Applied Ecology* 39.1 (Feb. 2002), pp. 8–17. ISSN: 0021-8901. DOI: 10.1046/j.1365-2664.2002.00704.x.

[3] C. Moritz et al. "Biogeographical concordance and efficiency of taxon indicators for establishing conservation priority in a tropical rainforest biota." In: *Proceedings. Biological sciences* 268.1479 (Sept. 2001), pp. 1875–1881. ISSN: 0962-8452. DOI: 10.1098/rspb.2001.1713.

[4] I Oliver and a J Beattie. "Invertebrate morphospecies as surrogates for species: a case study". In: *Conservation Biology* 10.1 (Feb. 1996), pp. 99–109. ISSN: 0888-8892. DOI: 10.1046/j.1523-1739.1996.10010099.x.

[5] Jiangning Wang et al. "A new automatic identification system of insect images at the order level". In: *Knowledge-Based Systems* 33 (Sept. 2012), pp. 102–110. ISSN: 09507051. DOI: 10.1016/j.knosys.2012.03.014.

[6] He-Ping Yang et al. "A tool for developing an automatic insect identification system based on wing outlines." In: *Scientific reports* 5.1 (Aug. 2015), p. 12786. ISSN: 2045-2322. DOI: 10.1038/srep12786.

[7] Paul J. D. Weeks and Kevin J. Gaston. "Image analysis, neural networks, and the taxonomic impediment to biodiversity studies". In: *Biodiversity and Conservation* 6.2 (1997), pp. 263–274. DOI: 10.1023/a:1018348204573.

[8] Seung-Ho Kang, Su-Hee Song, and Sang-Hee Lee. "Identification of butterfly species with a single neural network system". In: *Journal of Asia-Pacific Entomology* 15.3 (Sept. 2012), pp. 431–435. ISSN: 12268615. DOI: 10.1016/j.aspen.2012.03.006.

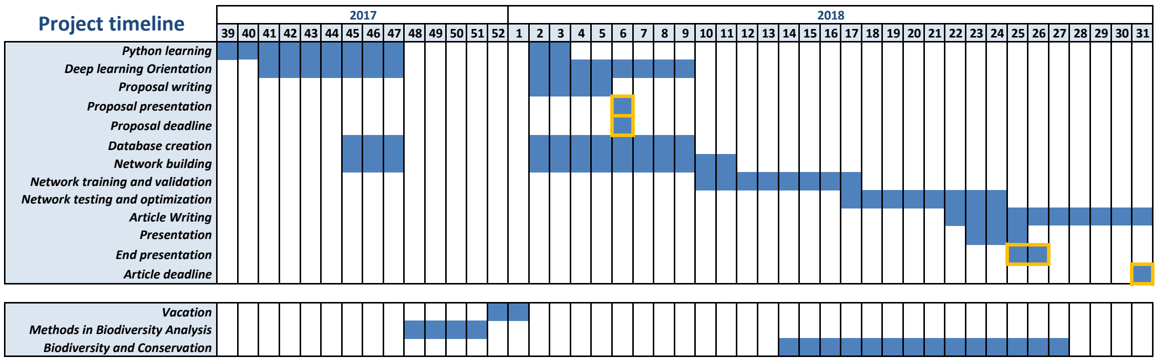

Figure 7. Timeline of the project. Orange boxed weeks represent deadlines. The bottom three parts of the timeline are not part of the actual project.

[9] Tiago Mauricio Francoy et al. "Identification of Africanized honey bees through wing morphometrics: two fast and efficient procedures". In: *Apidologie* 39.5 (Sept. 2008), pp. 488–494. ISSN: 0044-8435. DOI: 10.1051/apido:2008028.

[10] Maxence Gerard et al. "Discrimination of haploid and diploid males of Bombus terrestris (Hymenoptera; Apidae) based on wing shape". In: *Apidologie* 46.5 (Sept. 2015), pp. 644–653. ISSN: 0044-8435. DOI: 10.1007/s13592-015-0352-3.

[11] Neeraj Kumar et al. "Leafsnap: A Computer Vision System for Automatic Plant Species Identification". In: *Computer Vision – ECCV 2012*. 2012, pp. 502–516. ISBN: 978-3-642-33709-3. DOI: 10.1007/978-3-642-33709-3_36.

[12] Diah Harnoni Apriyanti, Aniati Murni Arymurthy, and Laksana Tri Handoko. "Identification of orchid species using content-based flower image retrieval". In: *2013 International Conference on Computer, Control, Informatics and Its Applications*. IEEE, Nov. 2013, pp. 53–57. ISBN: 9781479910786. DOI: 10.1109/IC3INA.2013.6819148. arXiv: 1406.2580.

[13] Chenglu Wen and Daniel Guyer. "Image-based orchard insect automated identification and classification method". In: *Computers and Electronics in Agriculture* 89 (Nov. 2012), pp. 110–115. ISSN: 01681699. DOI: 10.1016/j.compag.2012.08.008.

[14] Yann LeCun, Yoshua Bengio, and Geoffrey Hinton. "Deep learning". In: *Nature* 521.7553 (May 2015), pp. 436–444. ISSN: 0028-0836. DOI: 10.1038/nature14539. arXiv: 1603.05691.

[15] Yoshua Bengio, Aaron Courville, and Pascal Vincent. "Representation learning: A review and new perspectives". In: *IEEE Transactions on Pattern Analysis and Machine Intelligence* 35.8 (Aug. 2013), pp. 1798–1828. ISSN: 01628828. DOI: 10.1109/TPAMI.2013.50. arXiv: 1206.5538.

[16] David Held, Sebastian Thrun, and Silvio Savarese. "Deep Learning for Single-View Instance Recognition". In: *arXiv preprint arXiv:1507.08286* (July 2015). arXiv: 1507.08286.

[17] Dmytro Mishkin, Nikolay Sergievskiy, and Jiri Matas. "Systematic evaluation of CNN advances on the Imagenet". In: *arXiv preprint arXiv:1606.02228* 161 (Aug. 2017), pp. 11–19. DOI: 10.1016/j.cviu.2017.05.007. arXiv: 1606.02228.

[18] Sue Han Lee et al. "Deep-Plant: Plant Identification with convolutional neural networks". In: *arXiv preprint arXiv:1506.08425v1* (June 2015). ISSN: 15224880. DOI: 10.1109/ICIP.2015.7350839. arXiv: 1506.08425v1.

[19] Mads Dyrmann, Henrik Karstoft, and Henrik Skov Midtiby. "Plant species classification using deep convolutional neural network". In: *Biosystems Engineering* 151.2005 (Nov. 2016), pp. 72–80. ISSN: 15375110. DOI: 10.1016/j.biosystemseng.2016.08.024.

[20] Sharada P. Mohanty, David P. Hughes, and Marcel Salathé. "Using Deep Learning for Image-Based Plant Disease Detection". In: *Frontiers in Plant Science* 7.September (Sept. 2016), p. 1419. ISSN: 1664-462X. DOI: 10.3389/fpls.2016.01419. arXiv: 1604.03169.

[21] Hongwei Qin et al. "DeepFish: Accurate underwater live fish recognition with a deep architecture". In: *Neurocomputing* 187 (Apr. 2016), pp. 49–58. ISSN: 09252312. DOI: 10.1016/j.neucom.2015.10.122.

[22] Carlos Affonso et al. "Deep learning for biological image classification". In: *Expert Systems with Applications* 85 (Nov. 2017), pp. 114–122. ISSN: 09574174. DOI: 10.1016/j.eswa.2017.05.039.

[23] Gary Marcus. "Deep Learning: A Critical Appraisal". In: *arXiv preprint arXiv:1801.00631* (Jan. 2018), pp. 1–27. arXiv: 1801.00631.

[24] Hang Su et al. "Multi-view Convolutional Neural Networks for 3D Shape Recognition". In: *arXiv preprint arXiv:1505.00880* (May 2015). ISSN: 15505499. DOI: 10.1109/ICCV.2015.114. arXiv: 1505.00880.

[25] Charles R. Qi et al. "Volumetric and Multi-view CNNs for Object Classification on 3D Data". In: *arXiv preprint arXiv:1604.03265* (June 2016). ISSN: 10636919. DOI: 10.1109/CVPR.2016.609. arXiv: 1604.03265.

[26] Christian Peeters and Ross H Crozier. "Caste and Reproduction in Ants: Not All Mated Egg-Layers are "Queens"". In: *Psyche: A Journal of Entomology* 95.3-4 (1988), pp. 283–288. ISSN: 0033-2615. DOI: 10.1155/1988/52368.

[27] Peter Boer. "Are their native Holarctic Lasius and Serviformica ant species in the USA, other than exotic ones? With a key of the North American Lasius s.str. and Chthonolasius subgenera." English. Mar. 2016.

[28] AntWeb.org. *AntWeb*. URL: http://www.antweb.org (visited on 01/22/2017).

[29] AntWeb.org. *AntWeb API (version 2)*. URL: http://www.antweb.org/api/v2/ (visited on 01/22/2017).

[30] AntWeb.org. *AntWeb API (version 3)*. URL: https://www.antweb.org/api.do (visited on 01/22/2017).

[31] François Chollet et al. *Keras*. https://github.com/keras-team/keras. 2015.

[32] Martín Abadi et al. "TensorFlow: Large-Scale Machine Learning on Heterogeneous Distributed Systems". In: *arxiv preprint arXiv:1603.04467* (Mar. 2016). ISSN: 0270-6474. arXiv: 1603.04467.

[33] Kaiming He et al. "Deep Residual Learning for Image Recognition". In: *arXiv preprint arXiv:1512.03385* (Dec. 2015). ISSN: 15737721. arXiv: 1512.03385.

[34] Christian Szegedy et al. "Inception-v4, Inception-ResNet and the Impact of Residual Connections on Learning". In: *arXiv preprint arXiv:1602.07261* (Feb. 2016). arXiv: 1602.07261.

[35] Gao Huang et al. "Densely Connected Convolutional Networks". In: *arXiv preprint arXiv:608.06993* (Aug. 2016). ISSN: 0002-9645. arXiv: 1608.06993.

[36] Francois Chollet. "Xception: Deep Learning with Depthwise Separable Convolutions". In: *arXiv preprint arXiv:1610.02357*. IEEE, July 2017. ISBN: 978-1-5386-0457-1. DOI: 10.1109/CVPR.2017.195. arXiv: 1610.02357.

[37] Christian Szegedy et al. "Going Deeper with Convolutions". In: *arXiv preprint arXiv:1409.4842* (Sept. 2014). ISSN: 1749-4877. DOI: 10.1111/1749-4877.12092. arXiv: 1409.4842.

[38] Karen Simonyan and Andrew Zisserman. "Very Deep Convolutional Networks for Large-Scale Image Recognition". In: *arXiv preprint arXiv:1409.1556v6* (Sept. 2014). ISSN: 1095-9203. arXiv: 1409.1556v6.

[39] Nitish Srivastava et al. "Dropout: A Simple Way to Prevent Neural Networks from Overfitting". In: *Journal of Machine Learning Research* 15 (2014), pp. 1929–1958. ISSN: 15337928.

[40] Timothy Dozat. "Incorporating Nesterov Momentum into Adam". In: (2016).

[41] Diederik P. Kingma and Jimmy Ba. "Adam: A Method for Stochastic Optimization". In: *arXiv preprint arXiv:1412.6980* (Dec. 2014). ISSN: 09252312. arXiv: 1412.6980.

[42] Léon Bottou. "Stochastic Gradient Descent Tricks". In: *Neural networks: Tricks of the trade* (2012), pp. 421–436. ISSN: 03029743. DOI: 10.1007/978-3-642-35289-8. arXiv: arXiv:1206.5533v2.

[43] Ilya Loshchilov and Frank Hutter. "SGDR: Stochastic Gradient Descent with Warm Restarts". In: *arXiv preprint arXiv:1608.03983v5* (Aug. 2016), pp. 1–16. arXiv: 1608.03983v5.

[44] Sergey Ioffe and Christian Szegedy. "Batch Normalization: Accelerating Deep Network Training by Reducing Internal Covariate Shift". In: *arXiv preprint arXiv:1502.03167v3* (Feb. 2015). ISSN: 0717-6163. arXiv: 1502.03167v3.

[45] James Bergstra and Yoshua Bengio. "Random Search for Hyper-Parameter Optimization". In: *Journal of Machine Learning Research* 13.1 (2012), pp. 281–305. ISSN: 1532-4435.

[46] Matthew D Zeiler and Rob Fergus. "Visualizing and Understanding Convolutional Networks". In: *arXiv preprint arXiv:1311.2901v3* (Nov. 2014). ISSN: 978-3-319-10589-5. DOI: 10.1007/978-3-319-10590-1_53. arXiv: 1311.2901v3.

[47] Karen Simonyan, Andrea Vedaldi, and Andrew Zisserman. "Deep Inside Convolutional Networks: Visualising Image Classification Models and Saliency Maps". In: *arXiv preprint arXiv:1312.6034v2* (Dec. 2013). arXiv: 1312.6034v2.

[48] Andrej Karpathy, Justin Johnson, and Li Fei-Fei. "Visualizing and Understanding Recurrent Networks". In: *arXiv preprint arXiv:1506.02078v2* (June 2015). ISSN: 978-3-319-10589-5. arXiv: 1506.02078v2.